

# CONTAINERS AND KUBERNETES

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# DOCKER KNOWLEDGE

## Supplementary materials on Katacoda

<p>Deploying Your First Docker Container</p> <p>Learn how to launch containers using Docker</p> <p>Repeat Scenario</p>	<p>Deploy Static HTML Website as Container</p> <p>Learn how to run a static HTML website using Nginx</p> <p>Repeat Scenario</p>	<p>Building Container Images</p> <p>Learn how to build and launch your own container images</p> <p>Repeat Scenario</p>	<p>Dockerizing Node.js</p> <p>Learn how to deploy Node.js applications as containers</p> <p>Start Scenario</p>
<p>Optimise Builds With Docker OnBuild</p> <p>Learn how to optimise your Dockerfile using OnBuild</p> <p>Start Scenario</p>	<p>Ignoring Files During Build</p> <p>Learn how to ignore files being sent to the Docker Build Context or included in an image</p> <p>Start Scenario</p>	<p>Create Data Containers</p> <p>Learn how to use Data Containers and volumes-from property</p> <p>Start Scenario</p>	<p>Creating Networks Between Containers using Links</p> <p>Learn how containers communicate via links</p> <p>Repeat Scenario</p>
<p>Creating Networks Between Containers using Networks</p> <p>Learn how containers communicate via networks</p> <p>Start Scenario</p>	<p>Persisting Data Using Volumes</p> <p>Learn how to persist and share data between containers using Volumes</p> <p>Start Scenario</p>	<p>Manage Container Log Files</p> <p>Learn the different approaches for handling container log files</p> <p>Start Scenario</p>	<p>Ensuring Container Uptime With Restart Policies</p> <p>Understand how you can use Docker to ensure your containers stay up</p> <p>Start Scenario</p>

# CONTAINERS

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

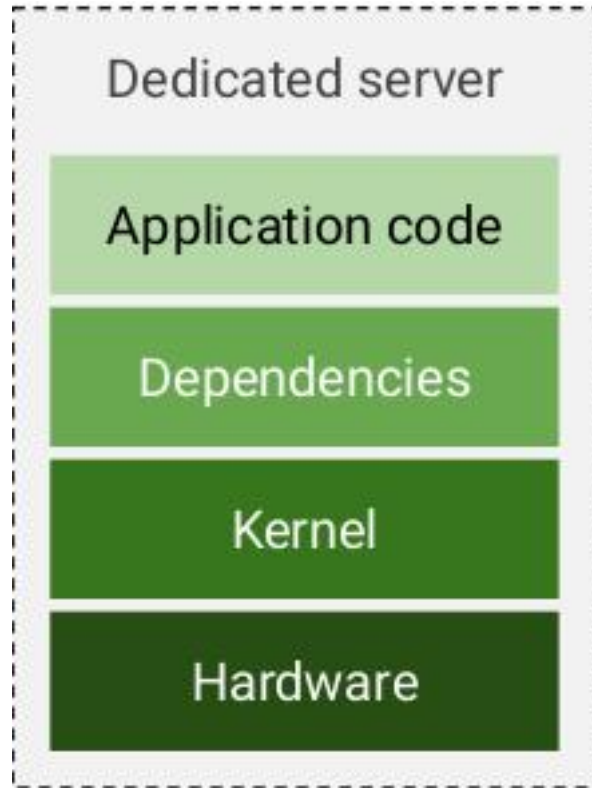
# PRE-REQUISITES

**In this section, we assume that you know enough about Docker and Kubernetes  
And, we will cover running these in the cloud**

**However,**

- To be on the same page,
- let's dedicate a few slides to a quick overview

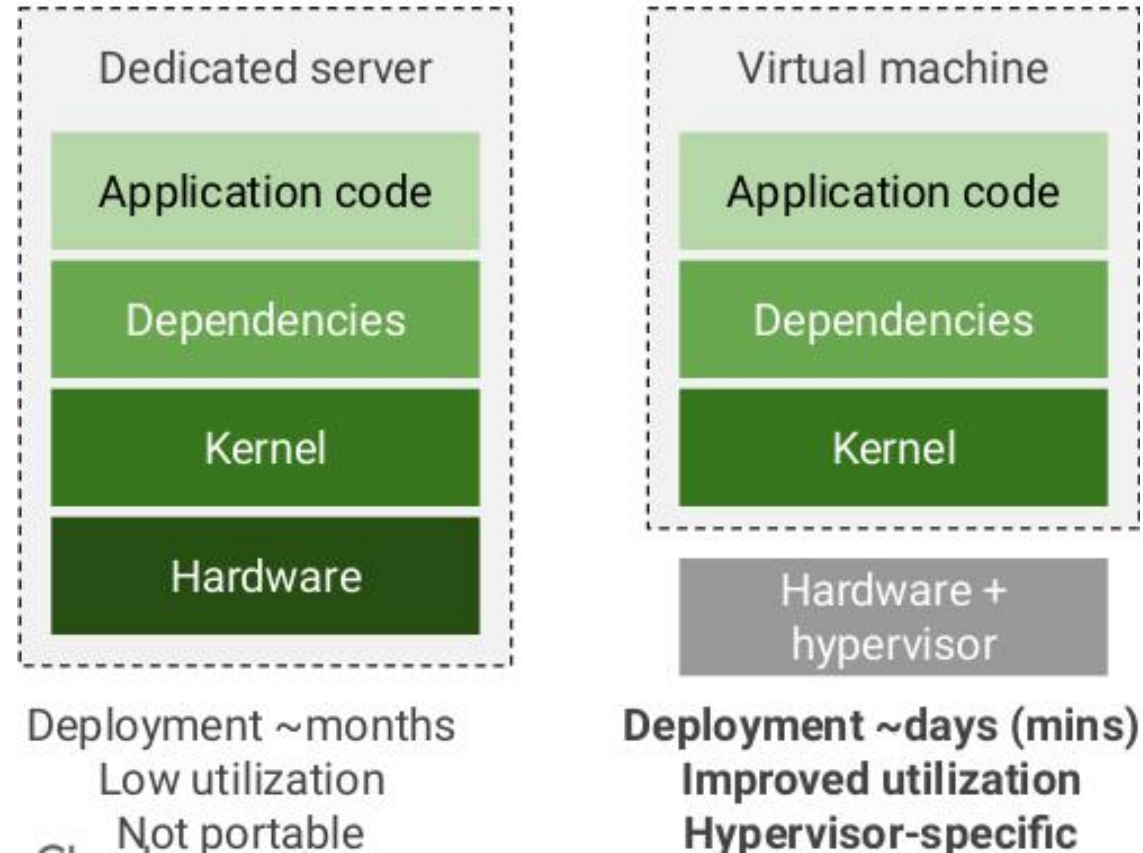
# BUILDING FROM SCRATCH



**Deployment ~months**  
**Low utilization**  
**Not portable**

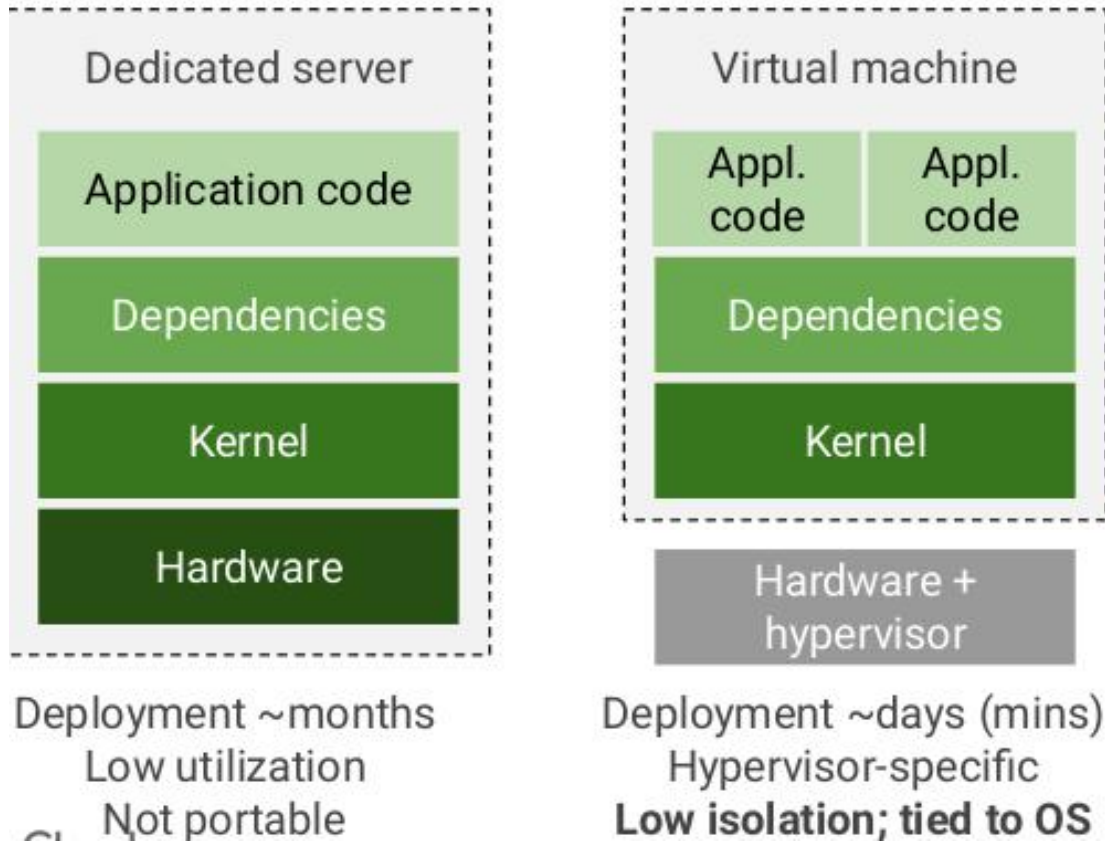
# VM STEP

Then VMware popularized running multiple servers and operating systems on the same hardware



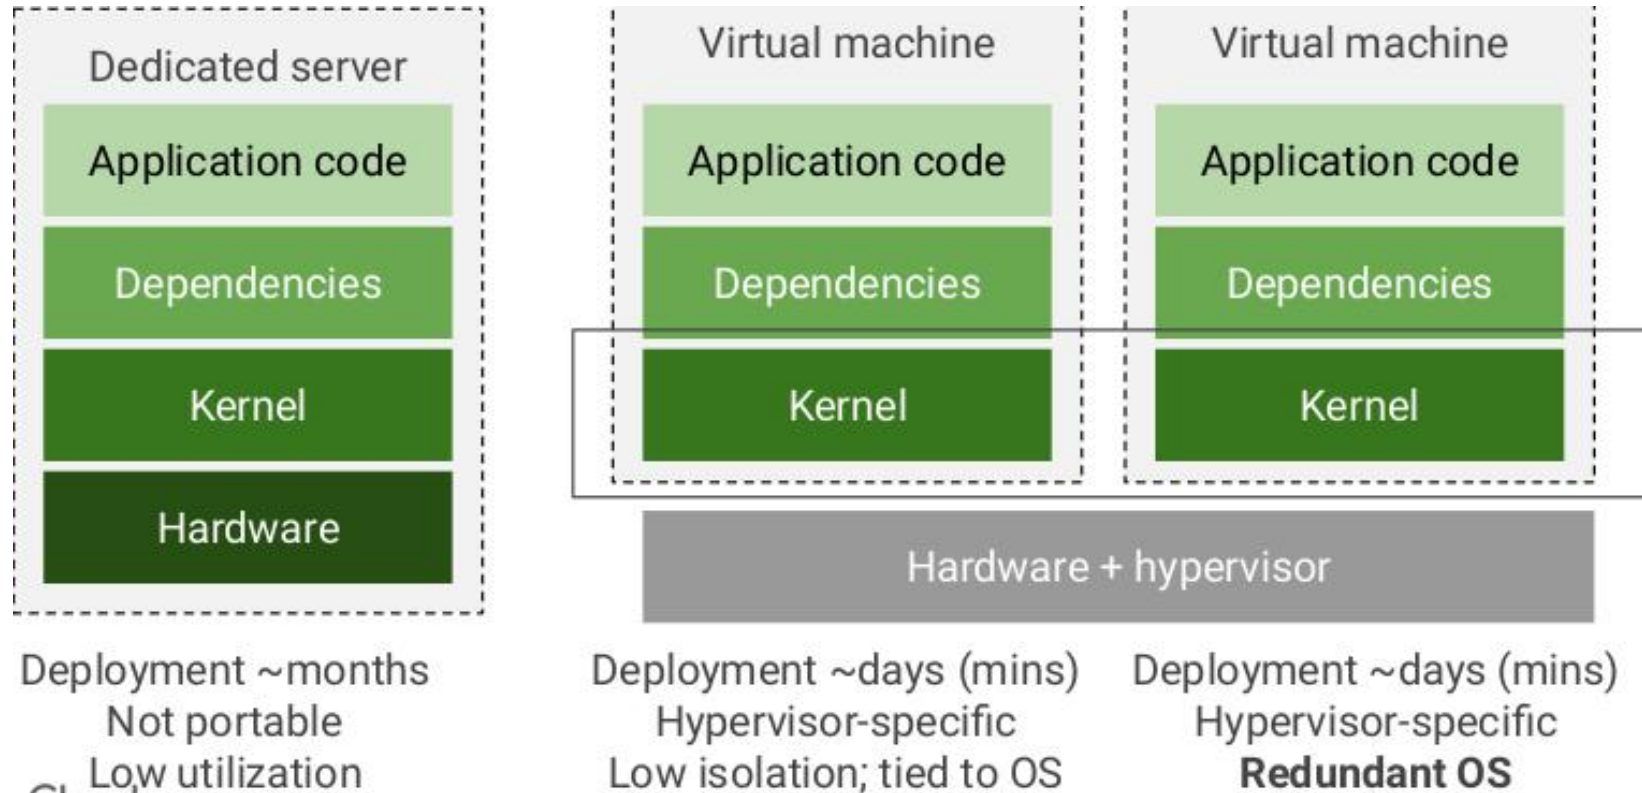
# STILL HARD

it was difficult to run and maintain multiple applications on a single VM, even with policies



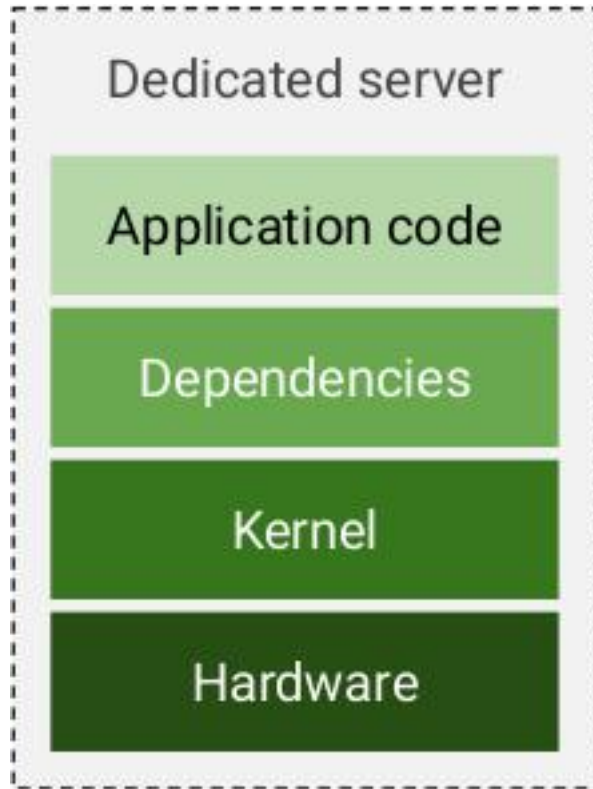
# STILL WASTEFUL

The VM-centric way to solve this is to run each app on its own server with its own dependencies, but that's wasteful

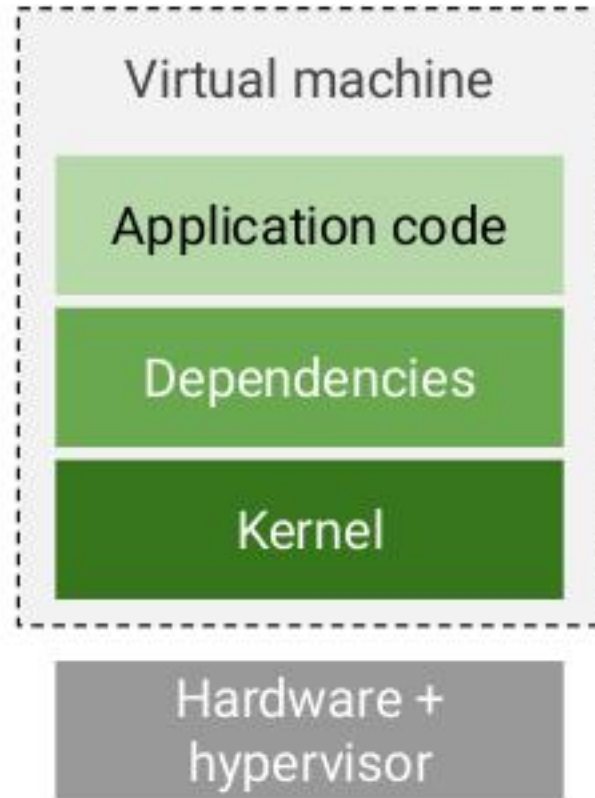




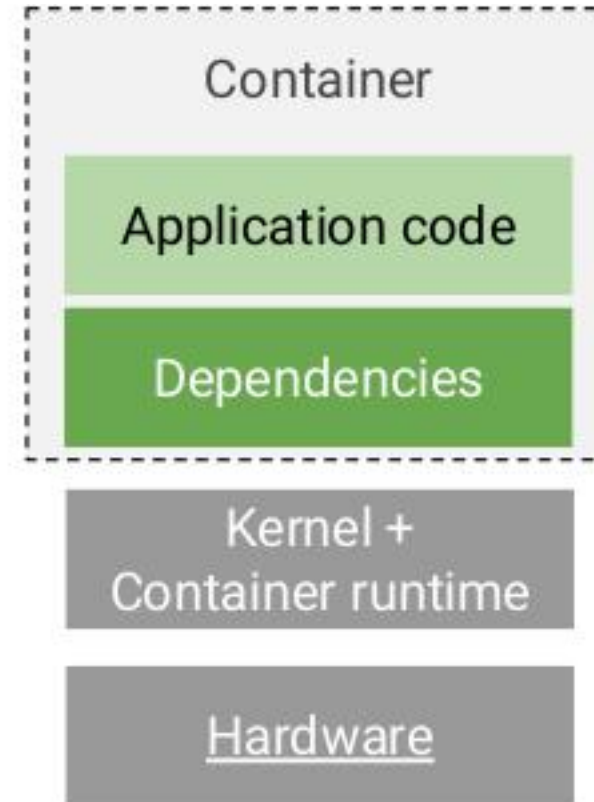
# VIRTUALIZE THE OS!



Deployment ~months  
Not portable  
Low utilization



Deployment ~days (mins)  
Hypervisor-specific  
Low isolation, Tied to OS



Deployment ~mins (sec)  
**Portable**  
**Very efficient**

# WHAT'S GOOD ABOUT CONTAINERS

## **Code works the same everywhere:**

- Across dev, test, & production
- Across bare-metal, VMs, cloud


## **Packaged apps speed development:**

- Agile creation and deployment
- Continuous integration/delivery
- Single file copy


## **They provide a path to microservices:**

- Introspectable
- isolated
- elastic

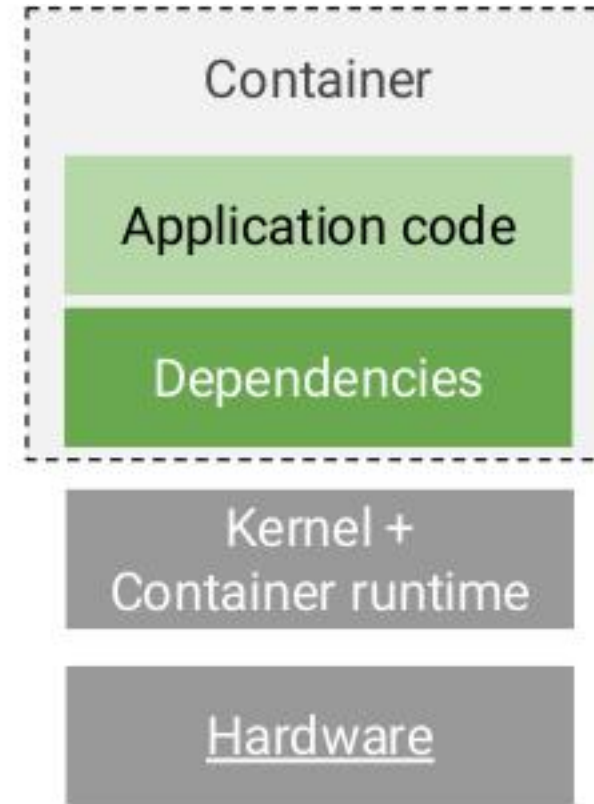
# VIRTUALIZE THE OS!



Deployment ~months  
Not portable  
Low utilization



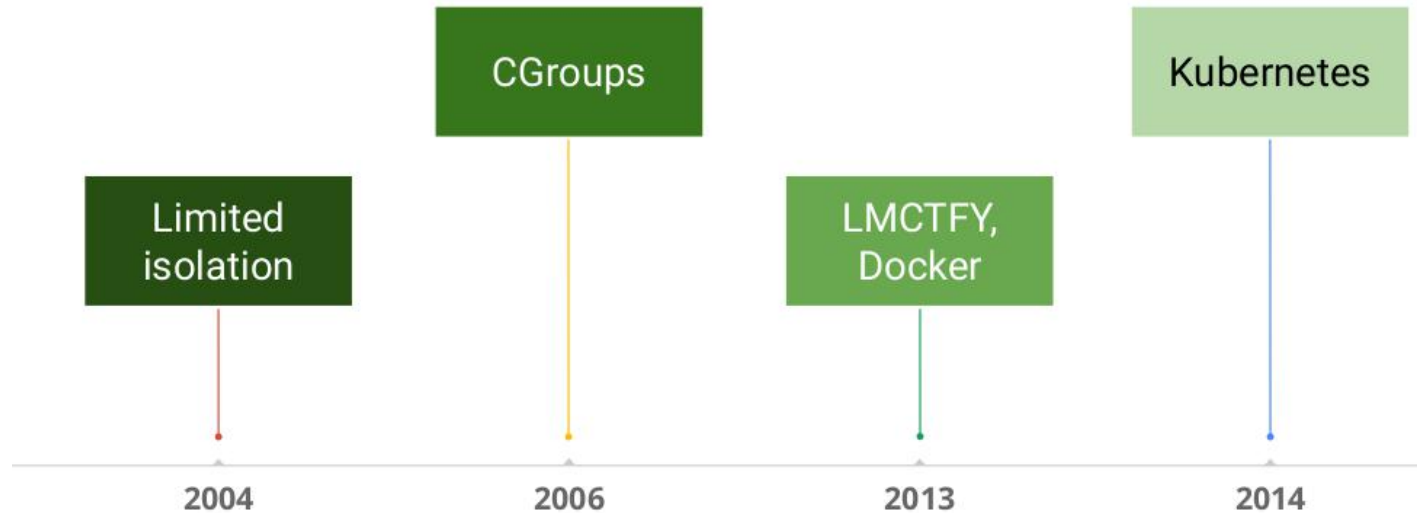
Deployment ~days (mins)  
Hypervisor-specific  
Low isolation, Tied to OS



Deployment ~mins (sec)  
**Portable**  
**Very efficient**

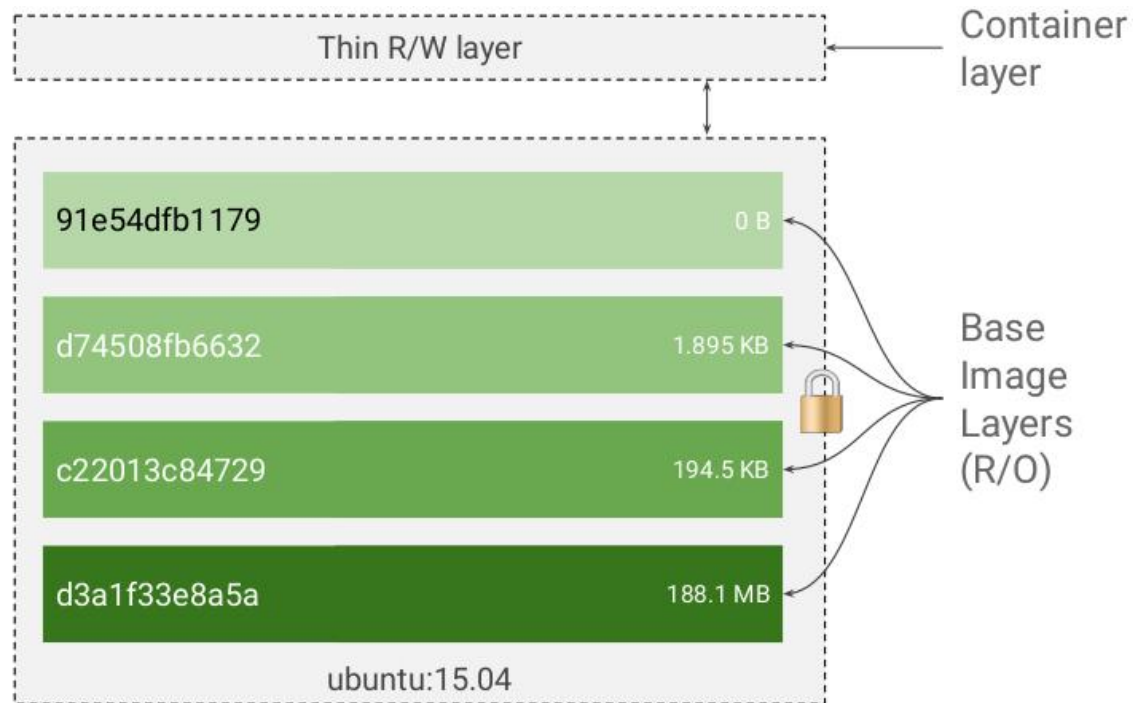
# CONTAINER HISTORY

Google has been developing and using containers to manage its applications for 12 years



# CONTAINERS' LAYERED FILE SYSTEM

Containers use a layered file system with only the top layer writable

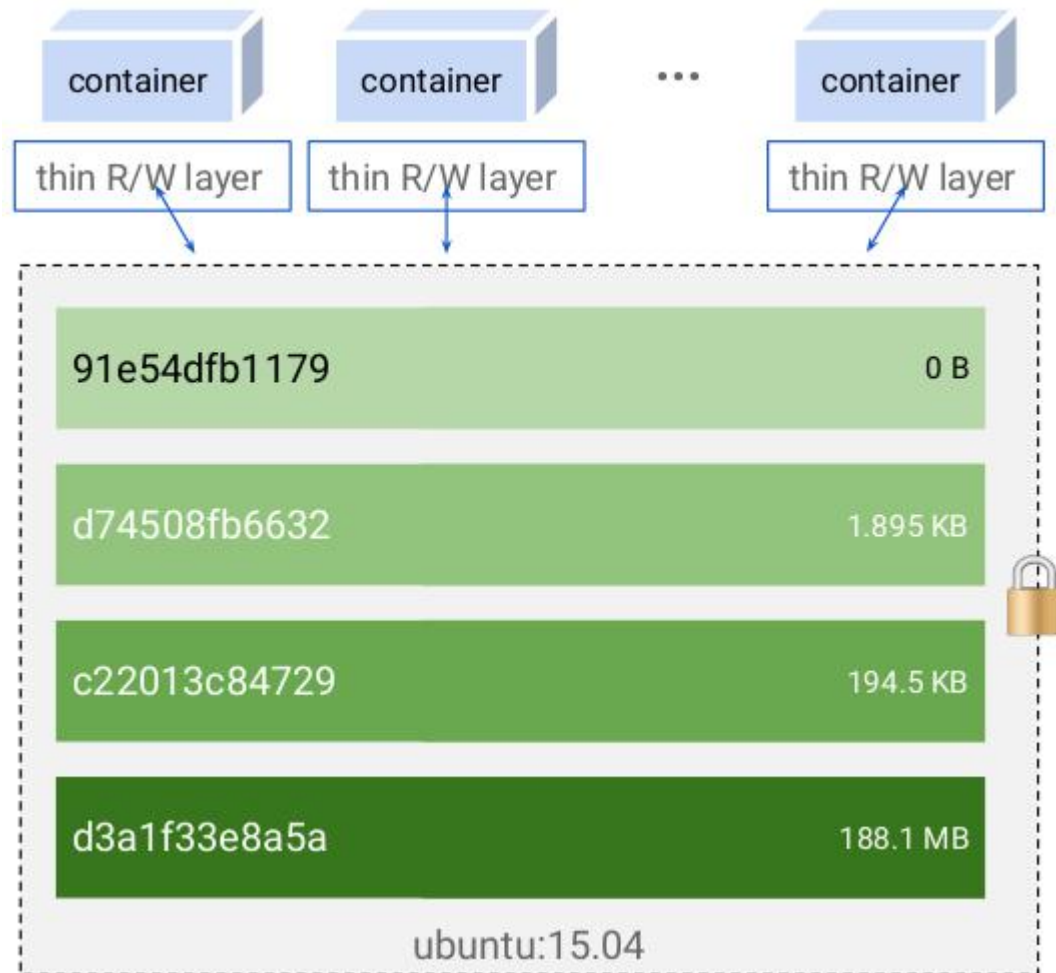


Dockerfile

```
FROM ubuntu:15.04
COPY . /app
RUN make /app
CMD python /app/app.py
```

# SMALLER SHARED IMAGES

Containers promote smaller shared images



# DOCKER

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**


**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# DOCKER ADOPTION



Source: Datadog

Source: DataDog



# DOCKERIZING EXAMPLE - CODE

This application listens on port 8080 and reply

```
python web-server.py
```

```
1 import tornado.ioloop
2 import tornado.web
3 import socket
4 dependencies
5 class MainHandler(tornado.web.RequestHandler):
6     def get(self):
7         self.write("Hostname: " +
8             socket.gethostname())
9     def make_app():
10         return tornado.web.Application([
11             (r"/", MainHandler),
12         ])
13 if __name__ == "__main__":
14     app = make_app()
15     listening on a port
16     app.listen(8080)
17     tornado.ioloop.IOLoop.current().start()
```

# CONTAINERIZE IT WITH DOCKER

```
$> docker build -t  
py-web-server .  
$> docker run -d py-web-server
```

```
FROM library/python:3.6.0-alpine  
RUN pip install tornado  
ADD web-server.py /web-server.py  
CMD ["python", "/web-server.py"]
```

You can also do stuff like:

```
$> docker images  
$> docker ps  
$> docker logs <container id>  
$> docker stop py-web-server
```

# USE GOOGLE DOCKER REGISTRY

```
docker build -t gcr.io/$PROJECT_ID/py-web-server:v1 .
```

— build a container image

```
gcloud docker -- push gcr.io/$PROJECT_ID/py-web-server:v1
```

— push it to a registry

```
docker run -d -p 8080:8080 --name py-web-server \
gcr.io/$PROJECT_ID/py-web-server:v1
```

— run it

# CONTAINERS ON GCP

**App Engine supports Docker containers as a custom runtime**

**Google Container Registry: private container image hosting on GCS with various CI/CD integrations**

**Compute Engine supports containers, including managed instance groups with Docker containers**

**The most powerful choice is a container orchestrator**



# QUIZ

**Docker Desktop is an app for building and sharing containerized apps and microservices available on which of the following operating systems?**

- A. macOS only
- B. Linux only
- C. Windows, macOS, and Windows Subsystem for Linux (WSL)

# QUIZ

**Which is correct Docker command to rebuild a container image?**

- A. `docker rebuild`
- B. `docker compile`
- C. `docker build`

# QUIZ

**Which of the following sentences describe a container image the best?**

- A. A container image is a read-only portable package that contains software and may include an operating system.
- B. A container image is a set of commands that builds a container.
- C. A container image is a read-only portable package that contains software.

# KUBERNETES

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

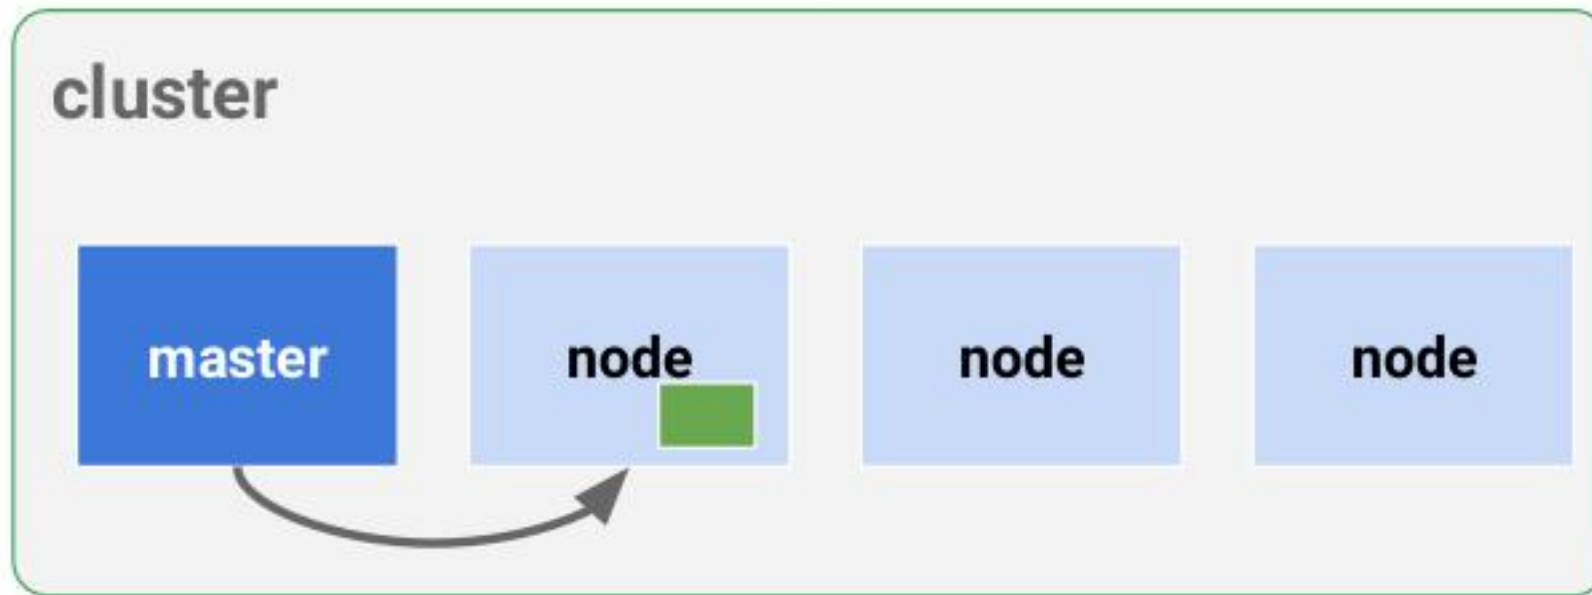
**KUBERNETES ON AZURE**



# WHAT IS KUBERNETES (K8)

Kubernetes manages jobs


Jobs run containers on nodes



# THE REAL WORLD

In a real ecosystem, a cluster may have 1000s of nodes and multiple masters.

Regional clusters have masters and nodes spread across 3 zones



# PARTICIPANT PRESENTATION

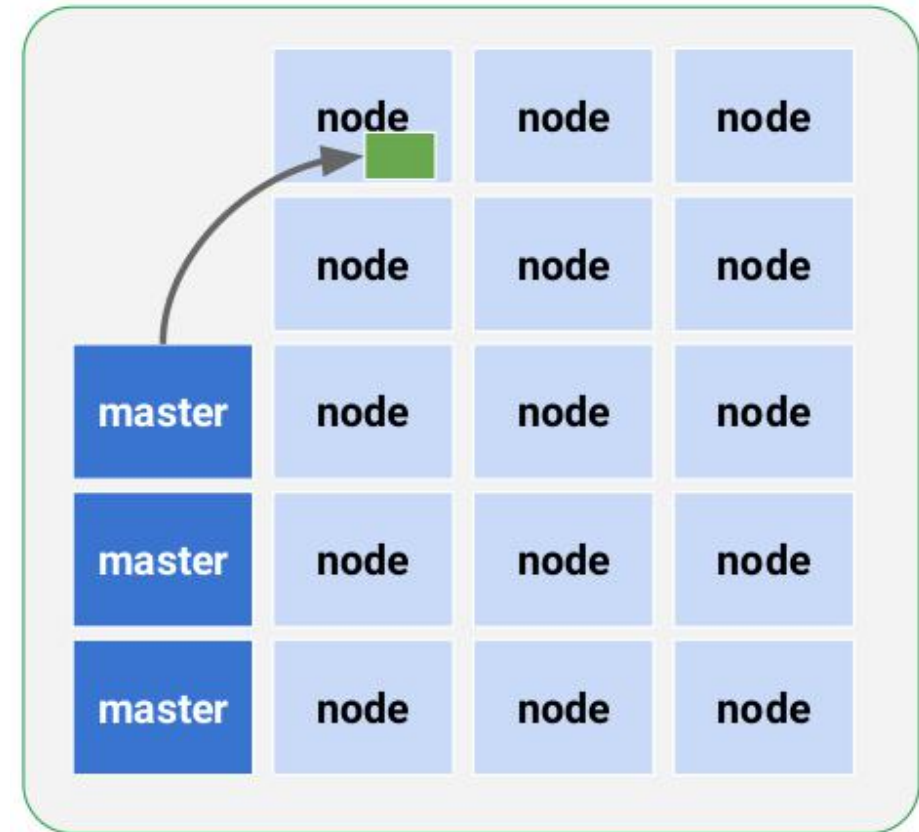
**Kubernetes use cases presented by the participants**

**Karl Kornel**

**Stanford team**

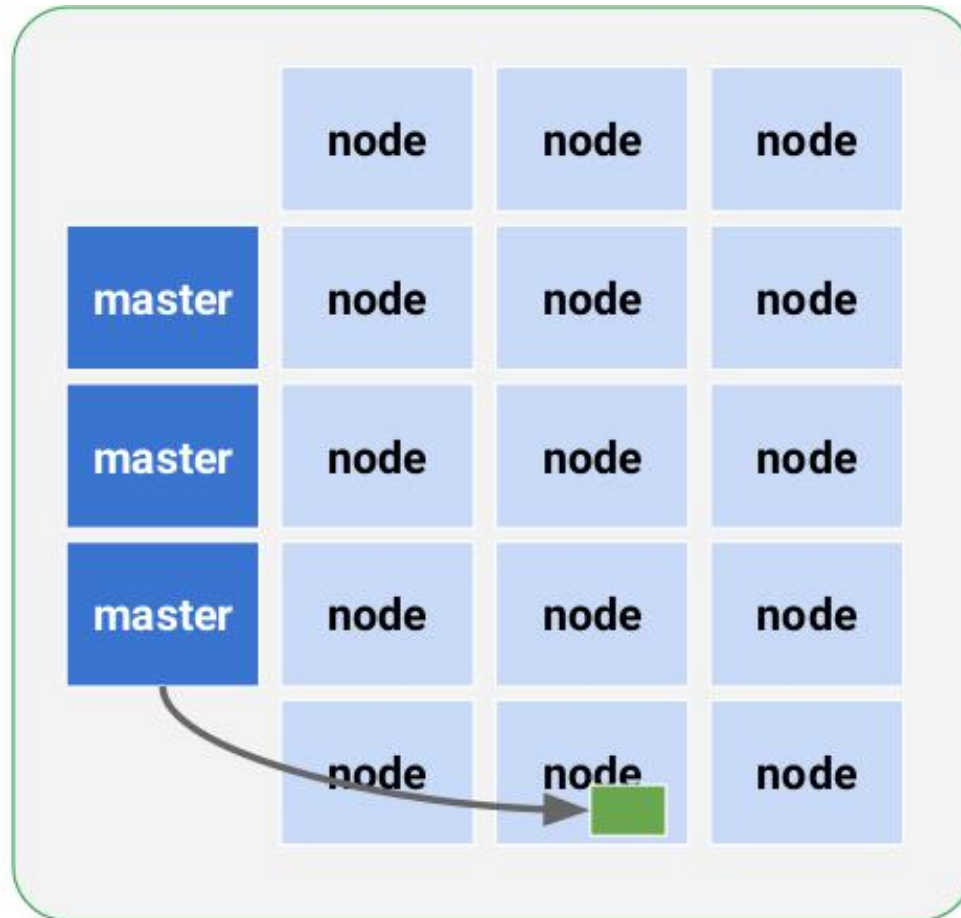
# MASTERS SCHEDULE JOBS

Masters schedule jobs on nodes based on load



# AND THEY CAN MOVE

And they can move them as needed to match resources




# POD

A job is called a Pod.

It's analogous to a VM

It can run multiple containers


It shares networking and storage separate from the node



# UNDER THE POD

Underneath the pod you have

- the node's hardware,
- OS, and
- NIC
- And a node IP




# INSIDE THE POD

Inside you have the pod with its containers

It gets its own network namespace with unique IP and set of ports

Data is stored in volumes in memory or persistent disks

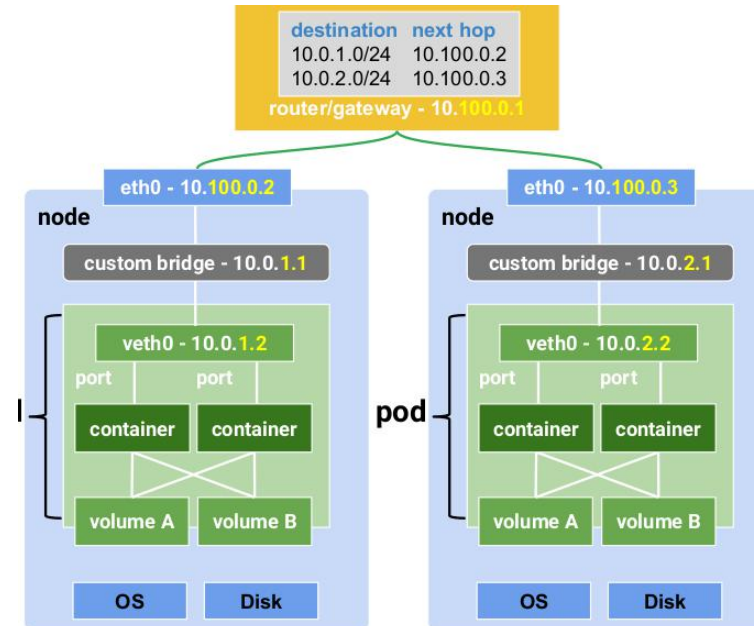




# GKE NETWORKING

Networking between pods in different nodes can be managed in a master routing table or other means.

GKE uses iptables in nodes and port forwarding for fewer points of failure (not shown).



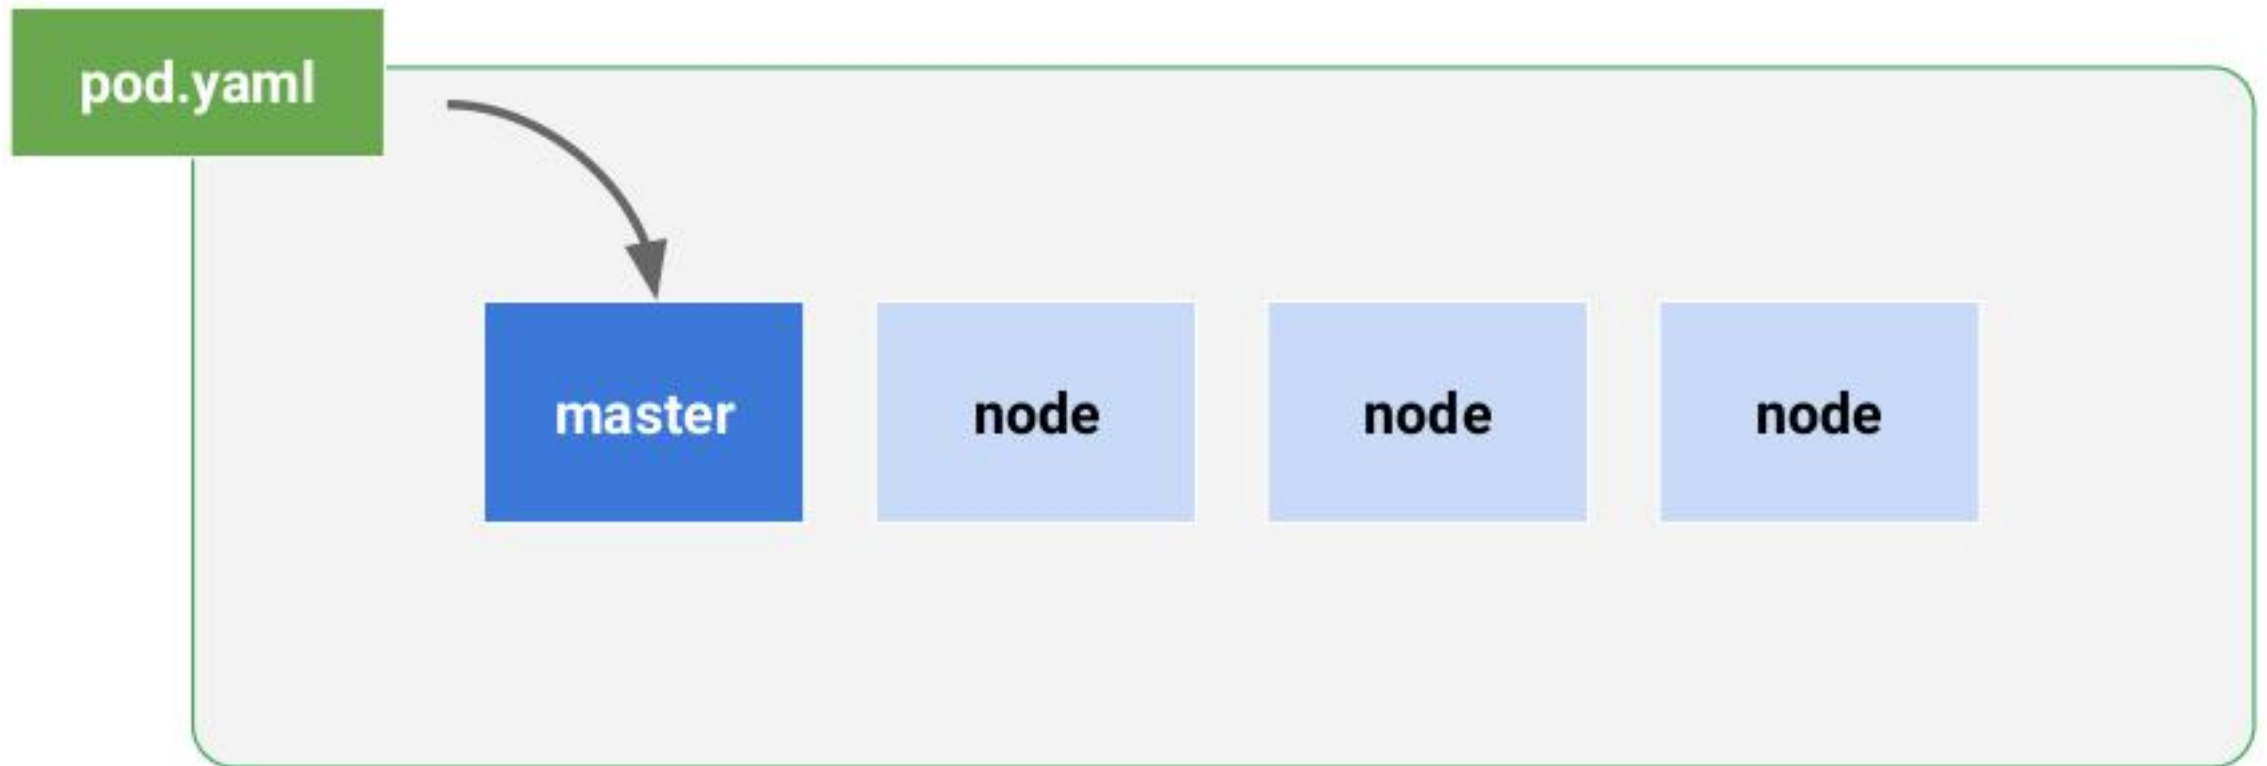
# POD IN YAML

You define a pod with a YAML file

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: my-app
5 spec:
6   containers:
7   - name: my-app
8     image: my-app
9   - name: nginx-ssl
10    image: nginx
11  ports:
12  - containerPort: 80
13  - containerPort: 443
```

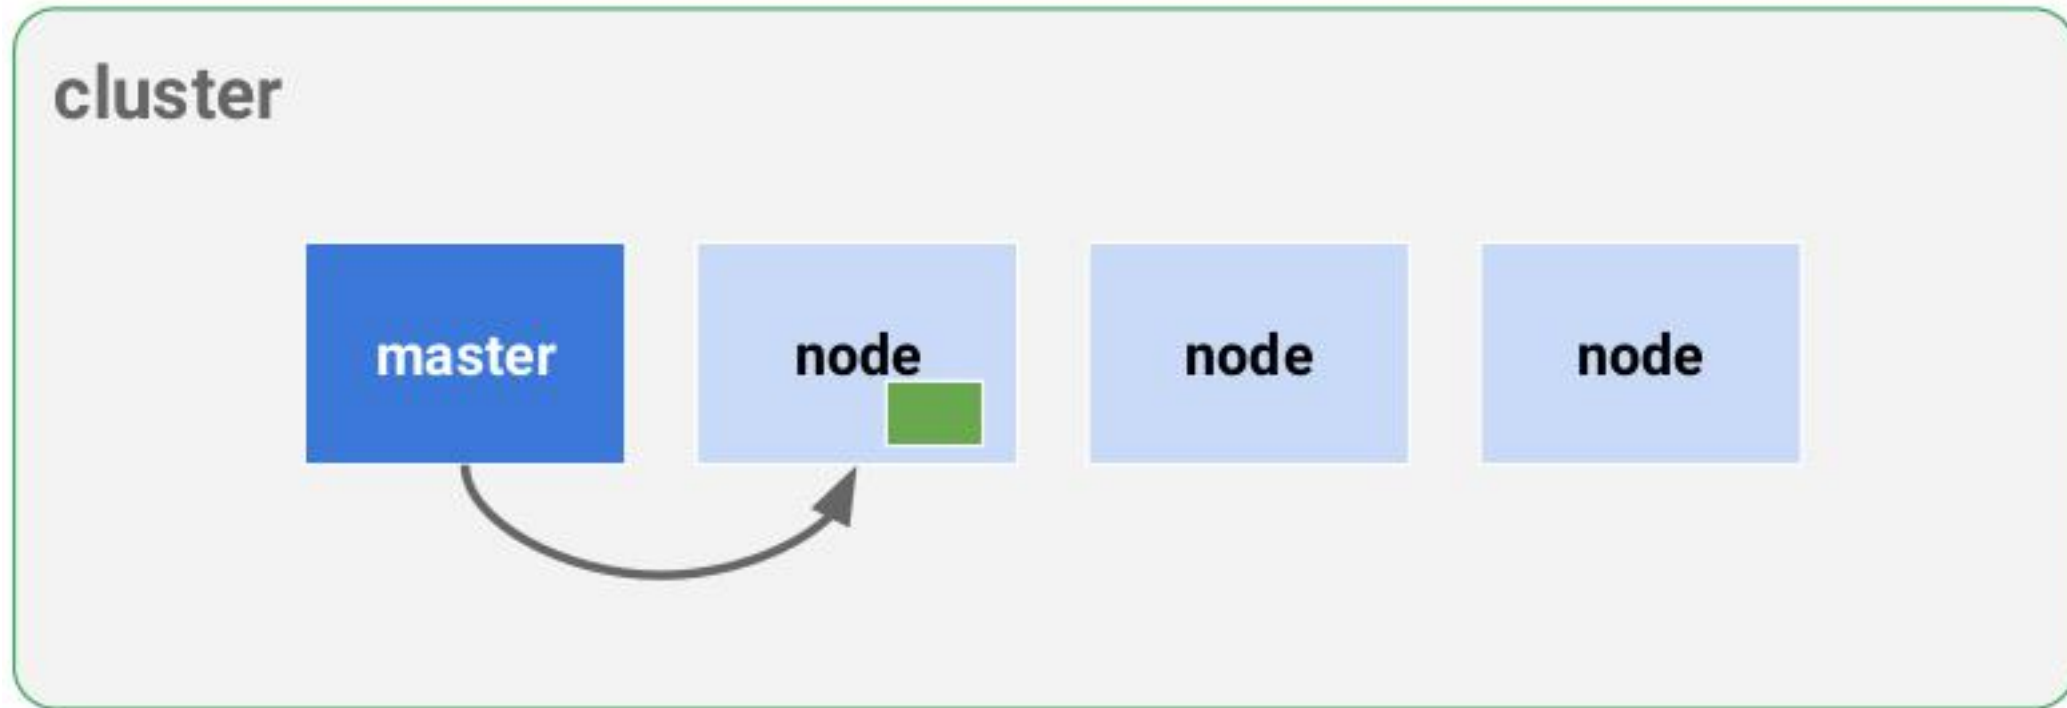
# YAML UPLOAD

You upload the YAML file to the master



# MASTER CREATES A POD

And the master creates a pod on your set of nodes



# POD COMPOSITION

A pod file is composed of several parts; for example...

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
  - name: my-app
    image: my-app
  - name: nginx-ssl
    image: nginx
  ports:
  - containerPort: 80
  - containerPort: 443
```

— API version

— pod resource

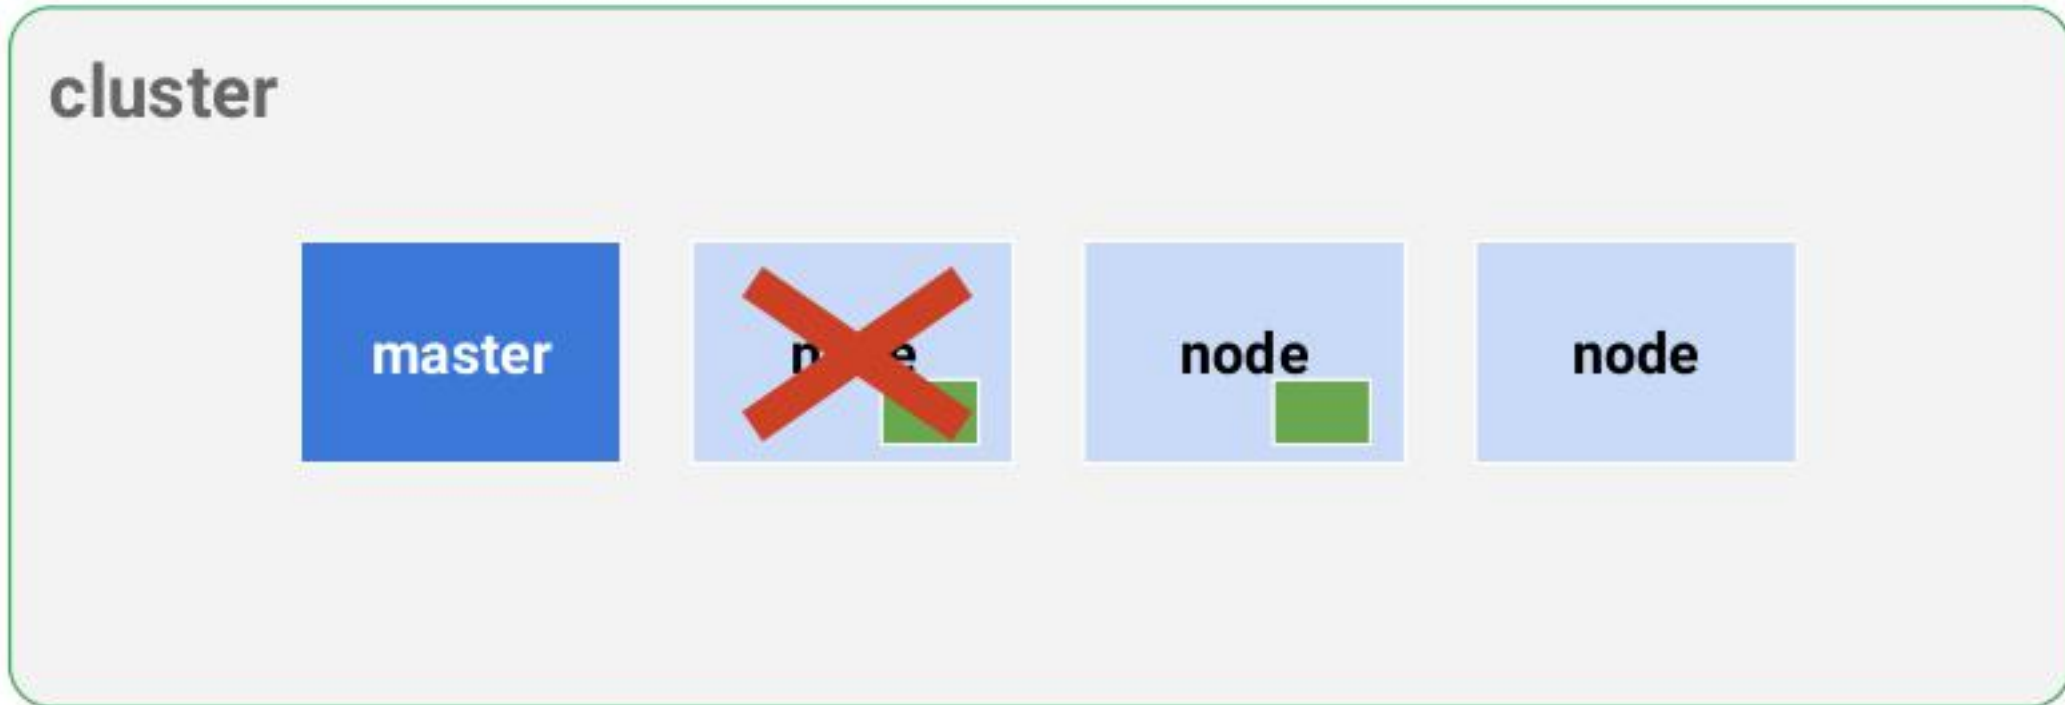
— pod name

— two containers

— NGINX front end on two ports

# MULTIPLE PODS

A deployment ensures that N pods are running in a cluster at any given time



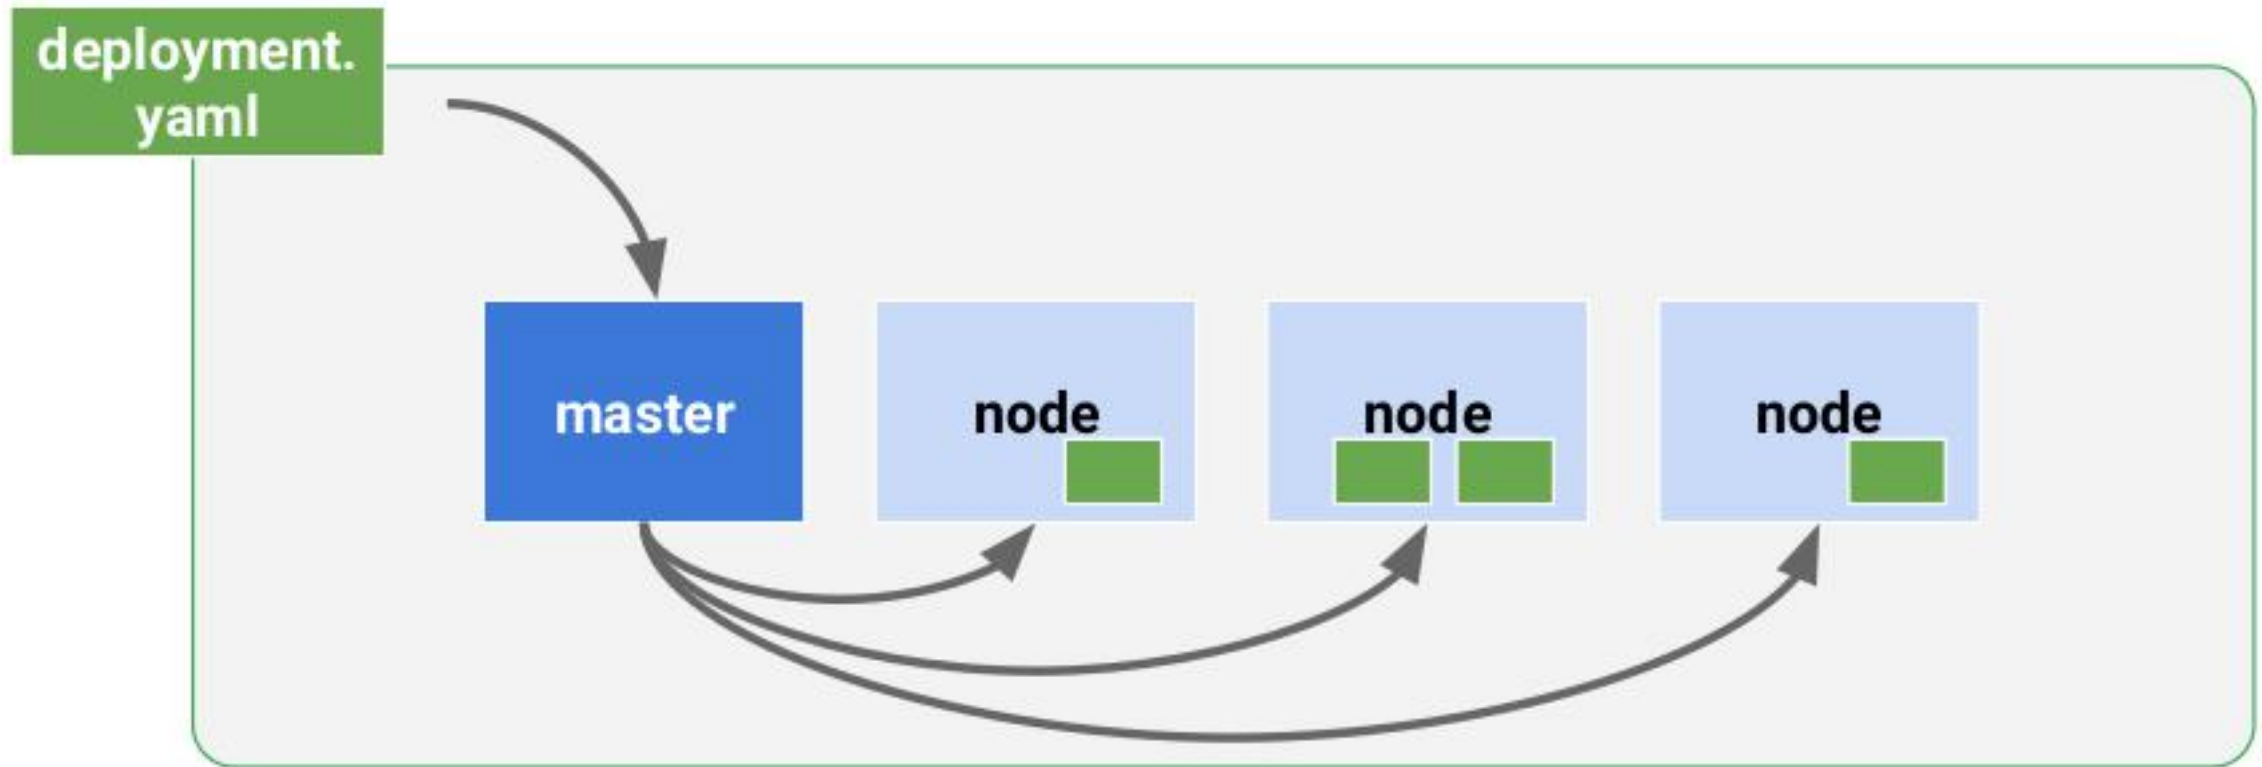
# YAML DEPLOYMENT

You define a deployment with a YAML file

```
kind: Deployment — deployment resource
apiVersion: v1.1
metadata:
  name: frontend — deployment name
spec:
  replicas: 4 — replicas
  selector: — pod selector
    role: web — role=web
  template:
    metadata:
      name: web
      labels: — pod label
        role: web — role=web
    spec:
      containers: — containers
      - name: my-app
        image: my-app
      - name: nginx-ssl
        image: nginx
        ports:
          - containerPort: 80
          - containerPort: 443
```

# RUNNING THE PODS

You upload the YAML file to the master, and the scheduler decides where to run the pods





# SERVICES, LABELS, SELECTORS

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**


**KUBERNETES ON AZURE**

# HOW PODS TALK TO EACH OTHER




# SERVICE

A service assigns a fixed IP to your pod replicas and allows other pods or services to communicate with them



# ENVIRONMENTS




# ENVIRONMENTS

Environment/Tier Name	Description
Environment/Tier Name	Description
Local	Developer's desktop/workstation
Development/Trunk	Development server acting as a sandbox where unit testing may be performed by the developer
Integration	CI build target, or for developer testing of side effects
Testing/Test/QC/Internal Acceptance	The environment where interface testing is performed. A quality control team ensures that the new code will not have any impact on the existing functionality and tests major functionalities of the system after deploying the new code in the test environment.
Staging/Stage/Model/Pre-production/External-Client Acceptance/Demo	Mirror of production environment
Production/Live	Serves end-users/clients

# MULTIPLE SERVICES

You can have multiple services with different configurations and features



To participants:

- Can you give application examples that would be relevant to your area of research?

# YAML, AGAIN

You define a service with a YAML file

```
kind: Service      — resource
apiVersion: v1
metadata:
  name: web-frontend
spec:
  ports:           — ports
  - name: http
    port: 80
    targetPort: 80
    protocol: TCP
  selector:        — pod selector
    role: web
  type: LoadBalancer — type is load balancer
```

# LABELS ARE METADATA

Labels are metadata you can assign to any API object and represent identity

They are

- The only grouping mechanism for pods
- Search by selectors





# EXAMPLE


This example has four pods and three labels



# QUERY FOR LABELS


You can query for labels that map to a value like the entire app

You can query for labels that map to a value like the entire app



# SEARCH WITH MULTIPLE LABELS

Or narrow your search with multiple labels like your app's frontend



**App = MyApp, Role = FE**

# SEARCH BACK-END

You can search your app's backend




**App = MyApp, Role = BE**




# SEARCH TEST PHASE

You can search your app's test phase




# SEARCH PRODUCTION RELEASE

You can search your app's production release



# KUBELET

Kubelet checks whether the pod is alive and healthy; if it gets a negative response or no reply...




## When to use

- In a decomposed compute model
- In a workflow chain where one piece fails for some reason

# POD RESTART


Kubelet restarts the pod





# KUBELET CONTINUES

And continues until it gets a healthy reply



# VOLUMES

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# KUBERNETES VOLUMES


Now let's look at how to get data into your pod and store data persistently using Kubernetes volumes.

Even though a pod isn't meant to be persistent, its data may be.

Docker also has a concept of volumes though it is somewhat looser and less managed than a Kubernetes volume.

# DOCKER VOLUMES

Docker provides data storage for containers, but volumes do not provide sharing between containers or lifecycle management



# DISCUSSION

**Do you remember different storage tiers might be considered here?**

**Price/performance tradeoff**

**Would you know where to look for the slide?**

# QUIZ

**Which App Engine environment runs your application in Docker containers on Google Compute Engine virtual machines?**

- A. App Engine Standard Environment
- B. App Engine Flexible Environment

# QUIZ

**Kubernetes Engine's cluster autoscaler automatically add or remove nodes from the cluster as required**

- A. False
- B. True

# QUIZ


**Care should be taken when creating node pools in multizone clusters because node pools are automatically replicated to those zones**

- A. False
- B. True



# KUBERNETES VOLUMES


Kubernetes volumes allow containers in pods to share data and be stateful



# A VOLUME IS A DIRECTORY

A volume is just a directory, and how it gets created depends on its type


```
1 $> kubectl create <volume>
```



# POD CONSUMES THAT DATA


Then you create a pod that consumes that data

```
1 $> kubectl create -f pod.yaml
```




# VOLUME ATTACHED TO THE POD

The volume is attached to the pod and made available to containers before they are brought online



# DATA ACCESS

Once the volume is attached, data can be mounted into a container's file system




# DISCUSSION

**Do you think it possible to have multiple volumes?**

- For example, with different performance characteristics) mounted in a single container?
- Fast data and slow data?


# CONTAINER READS DATA

Then the container is run and can get the mounted data



# VOLUME LIFECYCLE

Some volumes share the lifecycle of their pod





# KUBERNETES ON GCP - GKE

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# GCP COMPUTE AND PROCESSING OPTIONS

	Compute Engine	Container Engine ***	App Engine Standard and Flexible	Cloud Functions
Support language	Any	Any	Java, Python, Go, NodeJS	Triggers
Usage model	IaaS	IaaS, PaaS	PaaS	Microservices
	Server	Cluster	Autoscaling managed servers	Serverless
Primary use case	General workloads	Container workloads	Scalable web applications Mobile backend applications	Lightweight Event Actions

# GOOGLE CONTAINER ENGINE (GKE)

## Fully-managed service

- Kubernetes software maintained
- SLA

## Docker format containers

## Autoscaling (CPU or memory)

## Stackdriver logging and monitoring

## Cloud VPN integration

- Hybrid cloud and on premise solutions

## Cloud IAM integration




Google Kubernetes Engine

# CONTAINER CLUSTER


**Each node runs:**

- Docker runtime
- Kubelet agent
  - Manages scheduled Docker containers

**Network proxy**



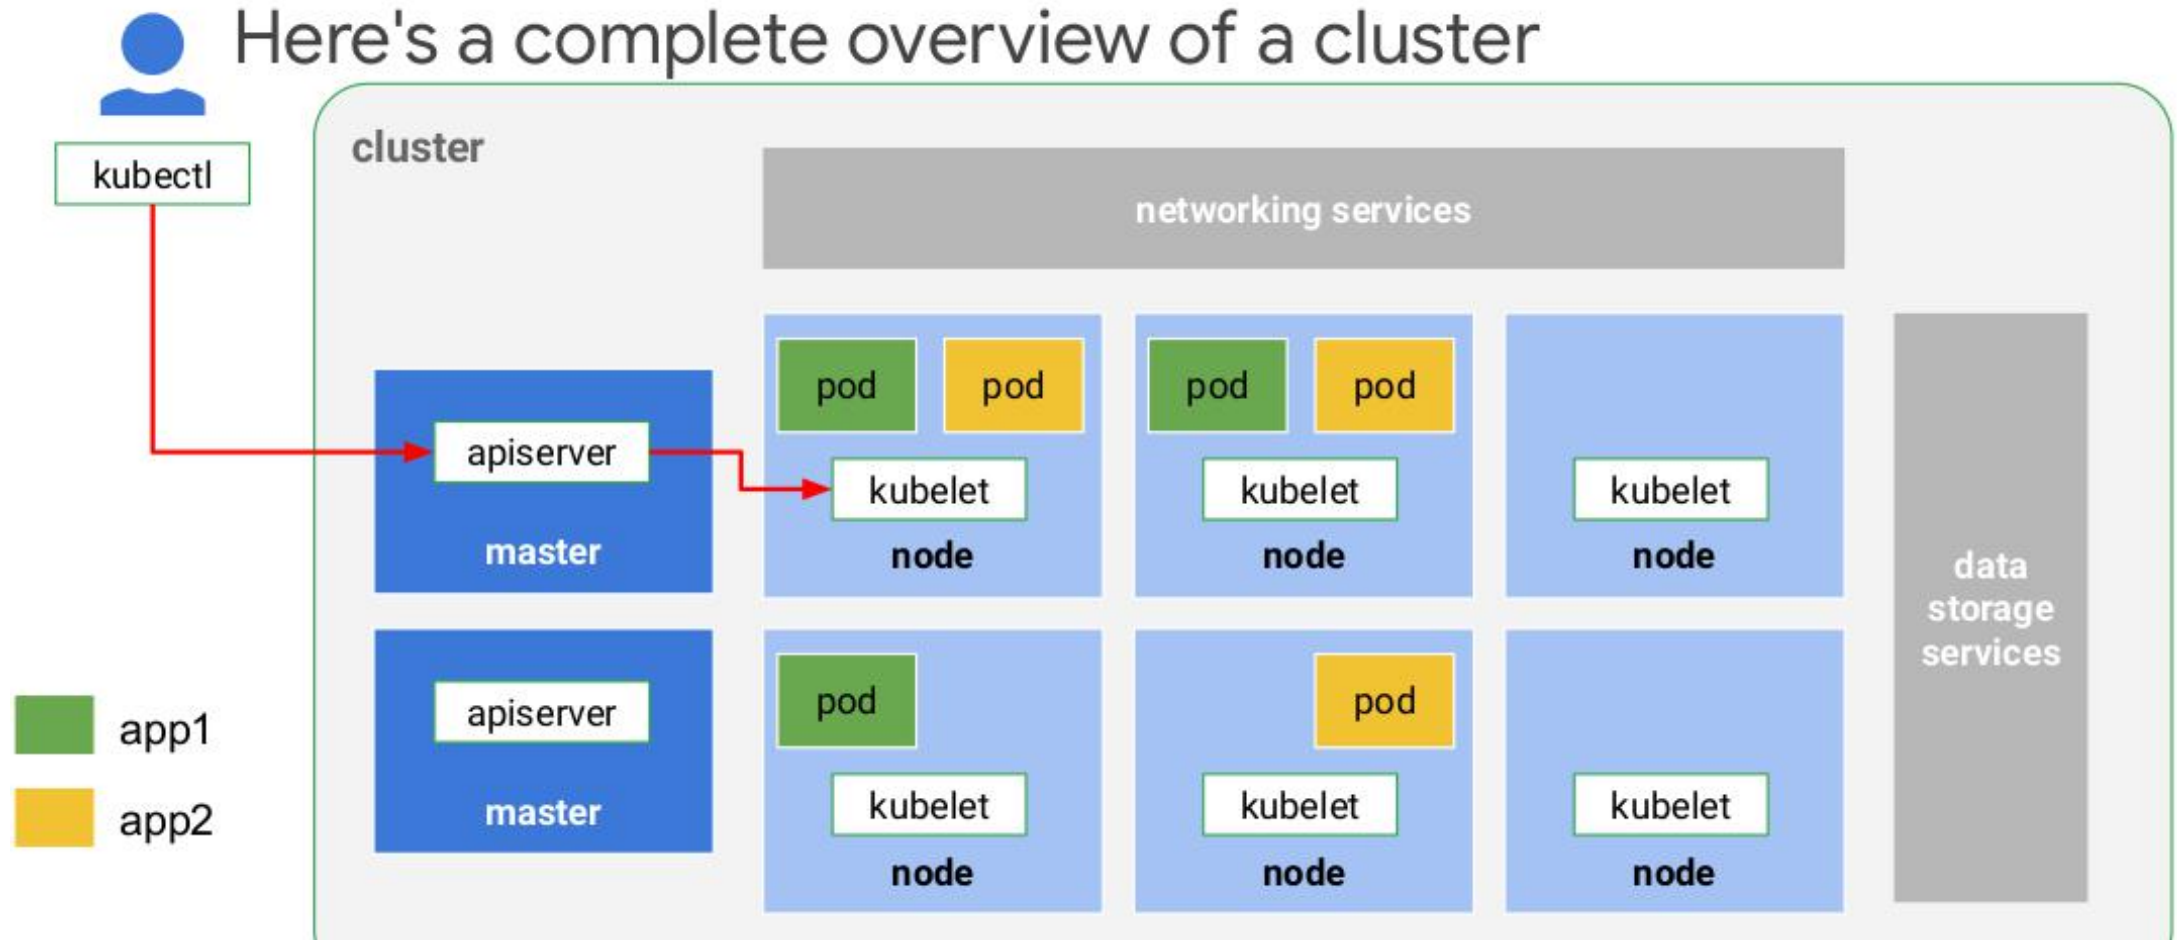
# KUBERNETES MASTER ENDPOINT



- Endpoint -- doorway to the cluster
- Kubernetes API server
  - Services REST requests
  - Schedules pod creation/deletion on nodes
  - Synchs pod info with service info
- Cloud Services integration



# CLUSTER VIEW



# KUBERNETES ON AWS

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# AMAZON ELASTIC CONTAINER REGISTRY

## Amazon Elastic Container Registry

Share and deploy container software, publicly or privately

[Get started with Amazon ECR](#)

[Browse public gallery](#)

### FEATURED

#### Amazon ECR Public - Share and Deploy Container Images Publicly

Share software worldwide for anyone to discover and download.




# RUNNING KUBERNETES ON AWS

## IF YOU WANT TO...

## CONSIDER USING

Fully manage your Kubernetes deployment. Provision and run Kubernetes on your choice of powerful instance types.

Amazon EC2



Run Kubernetes without needing to provision or manage master instances and etcd.

Amazon EKS



Store, encrypt, and manage container images for fast deployment.

Amazon ECR



# AMAZON EKS

## Amazon Elastic Kubernetes Service (Amazon EKS)

- Start, run, and scale Kubernetes applications in the AWS cloud or on-premises.
- Amazon EKS helps you provide highly-available and secure clusters
- Automates key tasks such as patching, node provisioning, and updates

## EKS runs upstream Kubernetes


- Certified Kubernetes conformant
- You can easily migrate any standard Kubernetes application to EKS without needing to refactor your code.

## Amazon EKS Distro

- Host and operate your Kubernetes clusters on-premises and at the edge
- AWS Outposts
- AWS Wavelength
- Amazon EKS Anywhere is coming in 2021.

# EKS WORKSHOP

multiple ways to configure VPC, ALB, and EC2 Kubernetes workers  
and Amazon Elastic Kubernetes Service



# CREATE USER

## Review

Review your choices. After you create the user, you can view and download the autogenerated password and access key.

## User details

User name	workshop
AWS access type	AWS Management Console access - with a password
Console password type	Custom
Require password reset	No
Permissions boundary	Permissions boundary is not set

## Permissions summary

The following policies will be attached to the user shown above.

Type	Name
Managed policy	<a href="#">AdministratorAccess</a>

# LOGIN TO AWS WORKSHOP PORTAL

The workshop creates an AWS account and a Cloud9 environment

## Login to AWS Workshop Portal

This workshop creates an AWS account and a Cloud9 environment. You will need the **Participant Hash** provided upon entry, and your email address to track your unique session.

Connect to the portal by clicking the button or browsing to <https://dashboard.eventengine.run/>. The following screen shows up.



Who are you?

### Terms & Conditions:

1. By using the Event Engine for the relevant event, you agree to the Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivate works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse

# LAUNCH CLOUD9

## Login to AWS Workshop Portal

This workshop creates an AWS account and a Cloud9 environment. You will need the **Participant Hash** provided upon entry, and your email address to track your unique session.

Connect to the portal by clicking the button or browsing to <https://dashboard.eventengine.run/>. The following screen shows up.



Who are you?

### Terms & Conditions:

1. By using the Event Engine for the relevant event, you agree to the Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivate works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse

# IN CLOUD9

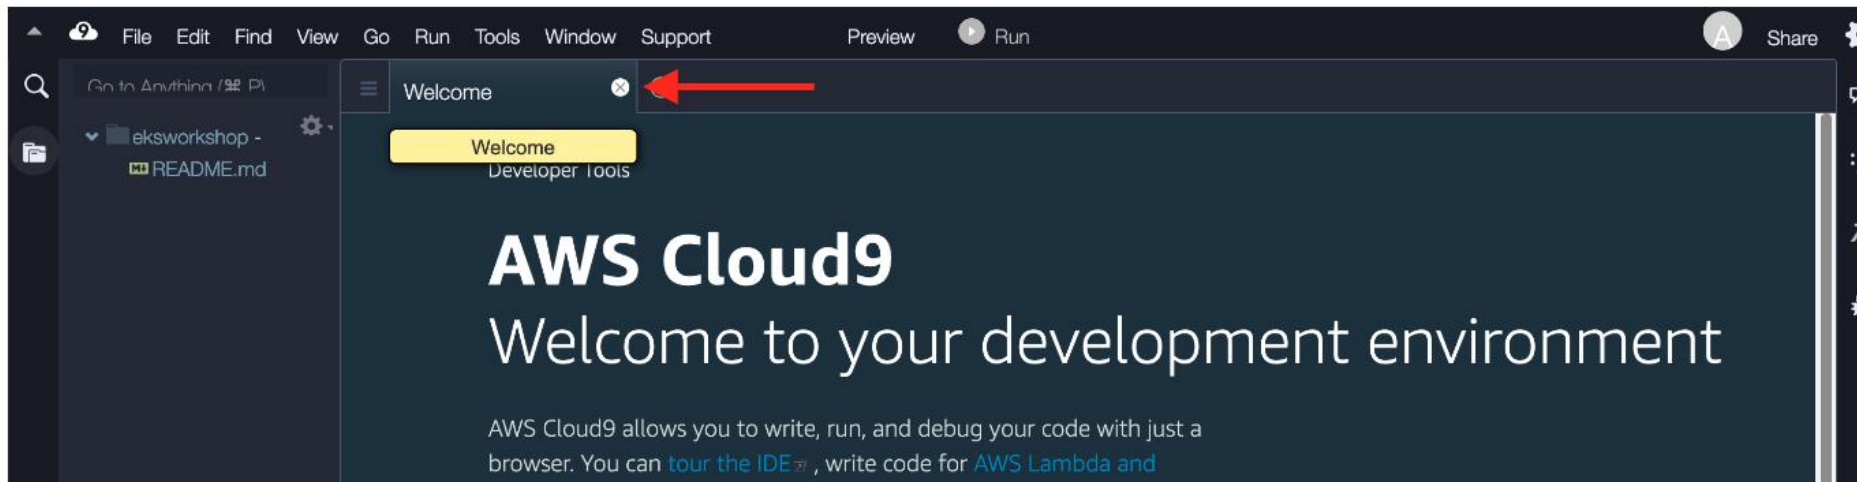
Oregon Ireland Ohio Singapore

Create a Cloud9 Environment: <https://us-west-2.console.aws.amazon.com/cloud9/home?region=us-west-2>

- Select **Create environment**
- Name it **eksworkshop**, click Next.
- Choose **t3.small** for Instance type, take all default values and click **Create environment**

When it comes up, customize the environment by:

- Closing the **Welcome tab**



# ON COMMAND-LINE

```
1 sudo curl --silent --location -o /usr/local/bin/kubectl \
2   https://amazon-eks.s3.us-west-2.amazonaws.com/1.17.11/2020-09-18/bin/linux/amd64/kubectl
3
4 sudo chmod +x /usr/local/bin/kubectl
```

**Install kubectl**

**Update awscli**

**Install jq, envsubst (from GNU gettext utilities) and bash-completion**

**Install yq for yaml processing**

**set the AWS Load Balancer Controller version**

**Create the role eksworkshop-admin**



# CREATE AN EKS CLUSTER

**Test the cluster**

```
1 eksctl create cluster -f eksworkshop.yaml
```

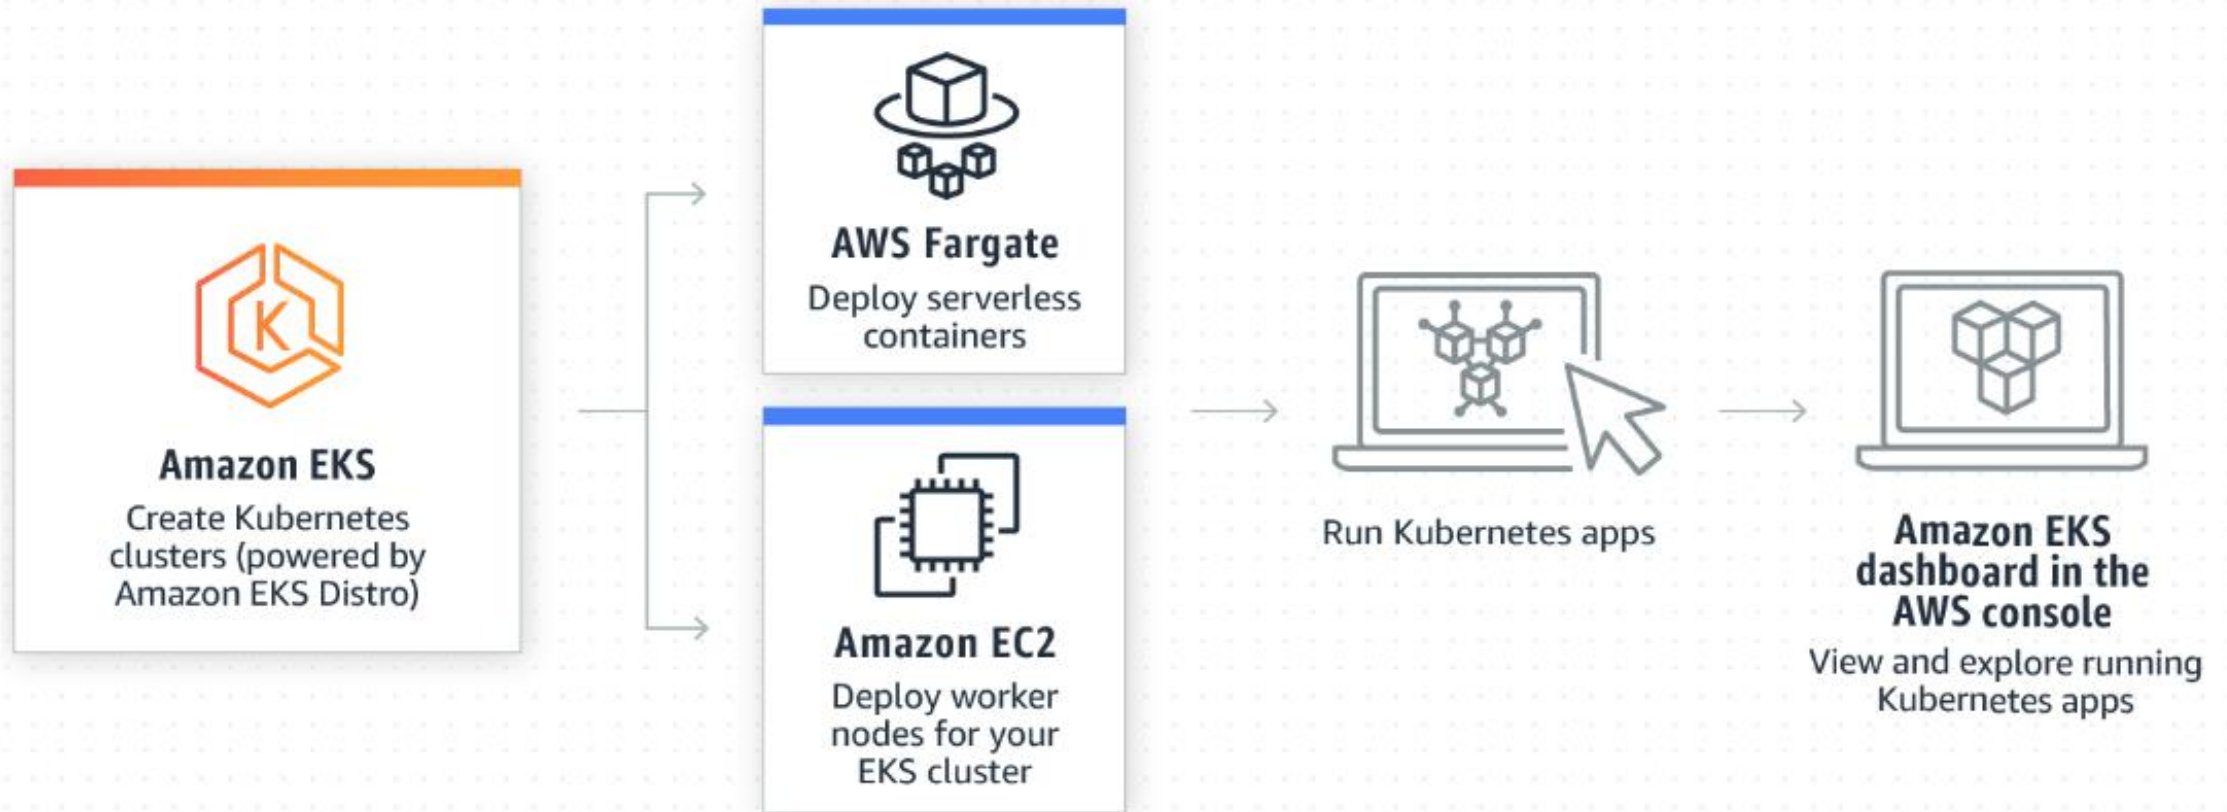
```
1 kubectl get nodes
```

**Export the Worker Role Name for use throughout the workshop**

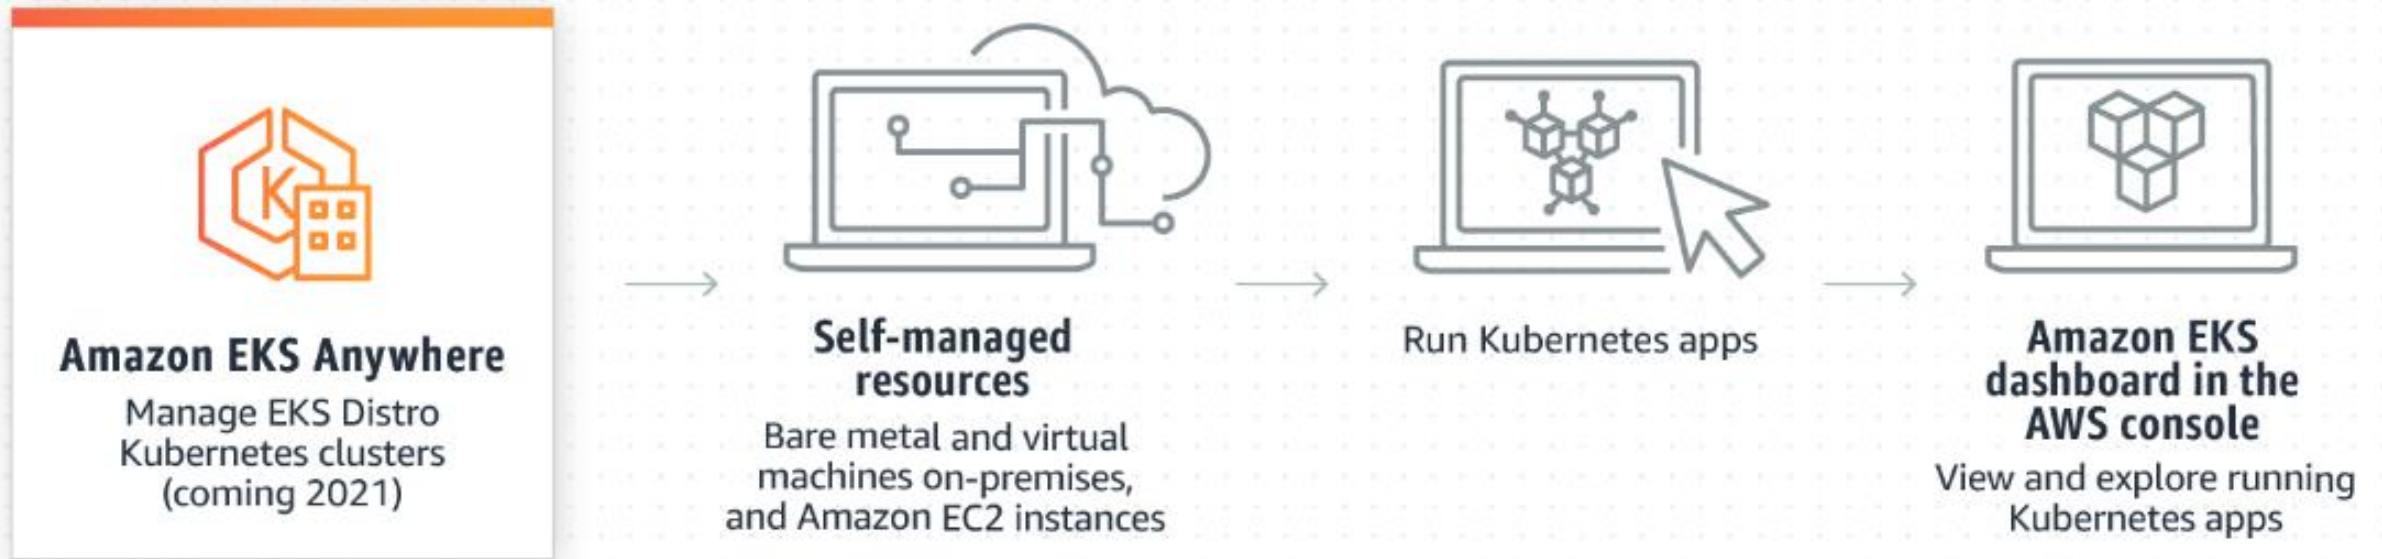
**Congratulations!**

**You now have a fully working Amazon EKS Cluster that is ready to use!**

# DEPLOY APPLICATIONS WITH AMAZON EKS



# DEPLOY APPLICATIONS WITH AMAZON EKS



# KUBERNETES ON AZURE

**CONTAINERS**

**DOCKER**

**KUBERNETES**

**SERVICES, LABELS, SELECTORS**

**VOLUMES**

**KUBERNETES ON GCP - GKE**

**KUBERNETES ON AWS**

**KUBERNETES ON AZURE**

# AZURE CONTAINER REGISTRY



The screenshot shows the Microsoft Azure website's navigation bar with the logo and links for Overview, Solutions, Products (highlighted with a dropdown arrow), Documentation, Pricing (with a dropdown arrow), and Training. Below the navigation bar is a breadcrumb trail: Home / Products / Container Registry. The main heading is 'Container Registry' in a large, bold font. Below it is a descriptive text: 'A registry of Docker and Open Container Initiative (OCI) images, with support for all OCI artifacts'. At the bottom of the section are two buttons: a solid blue 'Start free' button and a white 'Try now' button with a blue border.

Microsoft Azure

Overview Solutions **Products** Documentation Pricing Training

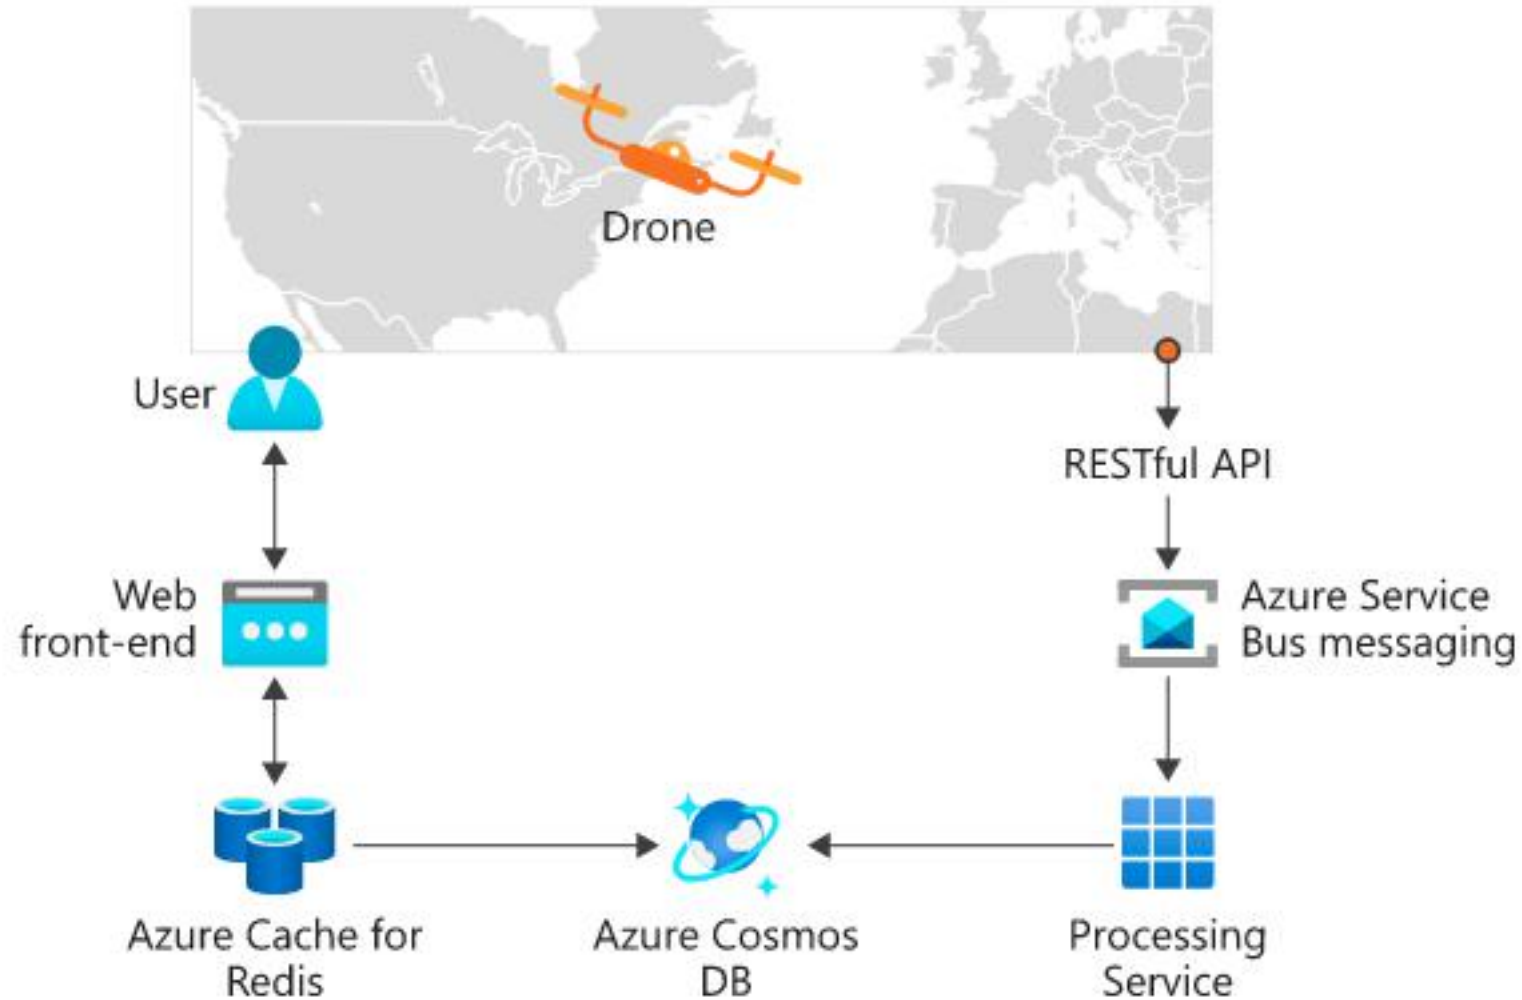
[Home](#) / [Products](#) / Container Registry

## Container Registry

A registry of Docker and Open Container Initiative (OCI) images, with support for all OCI artifacts

[Start free](#) [Try now](#)

# EXAMPLE OF KUBERNETES USE ON AZURE



# AZURE KUBERNETES SERVICE (AKS)

Elastic provisioning of capacity without the need to manage the infrastructure, and with the ability to add event-driven autoscaling and triggers through KEDA



# DISCUSSION

**Does anyone know of an example of autoscaling support from Research applications?**

**Outside of very big science?**

**(Sounds fairly fancy for a typical research project, but perhaps useful in sensor network applications, or where instruments can burst out data and events can be used to scale up data processing?)**



# AZURE KUBERNETES SERVICE (AKS)

**"Faster end-to-end development experience" (MSFT) through Visual Studio Code  
Kubernetes tools, Azure DevOps, and Azure Monitor**



# AZURE KUBERNETES SERVICE (AKS)

**Advanced identity and access management using Azure Active Directory, and dynamic rules enforcement across multiple clusters with Azure Policy**


- Single sign-on to any cloud app
- Enforce Multi-Factor Authentication with SaaS
- Works with multiple platforms and devices
- Integrate with on-premises Active Directory
- "Enterprise" (large production) Scale and SLA
- To discuss - can you think of relevant use cases?



# ANTHOS - MULTICLOUD FROM GOOGLE

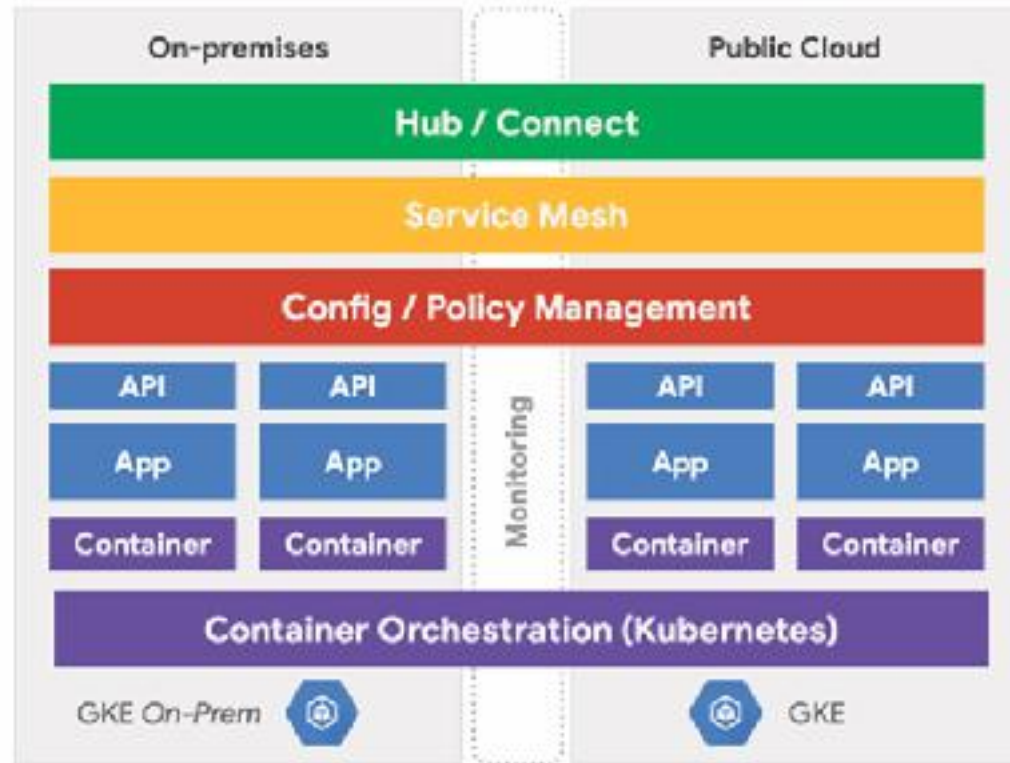


# WHAT IS ANTHOS?



(Image from ZDnet)

# ANTHOS AT A GLANCE



Anthos works on [AWS](#) and is in preview on Azure

Anthos will be covered in more detail in the next section

# CONGRATS ON COMPLETION

